

Forecasting: principles and practice

Rob J Hyndman

2.5 Seasonal ARIMA models

Outline

1 Backshift notation reviewed

2 Seasonal ARIMA models

3 ARIMA vs ETS

4 Lab session 12

Backshift notation

A very useful notational device is the backward shift operator, B , which is used as follows:

$$By_t = y_{t-1} .$$

Backshift notation

A very useful notational device is the backward shift operator, B , which is used as follows:

$$By_t = y_{t-1} .$$

In other words, B , operating on y_t , has the effect of **shifting the data back one period.**

Backshift notation

A very useful notational device is the backward shift operator, B , which is used as follows:

$$By_t = y_{t-1} .$$

In other words, B , operating on y_t , has the effect of **shifting the data back one period**. Two applications of B to y_t **shifts the data back two periods**:

$$B(By_t) = B^2y_t = y_{t-2} .$$

Backshift notation

A very useful notational device is the backward shift operator, B , which is used as follows:

$$By_t = y_{t-1} .$$

In other words, B , operating on y_t , has the effect of **shifting the data back one period**. Two applications of B to y_t **shifts the data back two periods**:

$$B(By_t) = B^2y_t = y_{t-2} .$$

For monthly data, if we wish to shift attention to “the same month last year,” then B^{12} is used, and the notation is $B^{12}y_t = y_{t-12}$.

Backshift notation

- First difference: $1 - B$.
- Double difference: $(1 - B)^2$.
- d th-order difference: $(1 - B)^d y_t$.
- Seasonal difference: $1 - B^m$.
- Seasonal difference followed by a first difference:
 $(1 - B)(1 - B^m)$.
- Multiply terms together to see the combined effect:

$$\begin{aligned}(1 - B)(1 - B^m)y_t &= (1 - B - B^m + B^{m+1})y_t \\ &= y_t - y_{t-1} - y_{t-m} + y_{t-m-1}.\end{aligned}$$

Backshift notation for ARIMA

■ ARMA model:

$$\begin{aligned}y_t &= c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} + e_t + \theta_1 e_{t-1} + \cdots + \theta_q e_{t-q} \\&= c + \phi_1 B y_t + \cdots + \phi_p B^p y_t + e_t + \theta_1 B e_t + \cdots + \theta_q B^q e_t\end{aligned}$$

$$\phi(B)y_t = c + \theta(B)e_t$$

$$\text{where } \phi(B) = 1 - \phi_1 B - \cdots - \phi_p B^p$$

$$\text{and } \theta(B) = 1 + \theta_1 B + \cdots + \theta_q B^q.$$

Backshift notation for ARIMA

■ ARMA model:

$$\begin{aligned}y_t &= c + \phi_1 y_{t-1} + \cdots + \phi_p y_{t-p} + e_t + \theta_1 e_{t-1} + \cdots + \theta_q e_{t-q} \\&= c + \phi_1 B y_t + \cdots + \phi_p B^p y_t + e_t + \theta_1 B e_t + \cdots + \theta_q B^q e_t\end{aligned}$$

$$\phi(B)y_t = c + \theta(B)e_t$$

$$\text{where } \phi(B) = 1 - \phi_1 B - \cdots - \phi_p B^p$$

$$\text{and } \theta(B) = 1 + \theta_1 B + \cdots + \theta_q B^q.$$

■ ARIMA(1,1,1) model:

$$\begin{array}{ccccc}(1 - \phi_1 B) & (1 - B)y_t & = & c + (1 + \theta_1 B)e_t \\ \uparrow & \uparrow & & \uparrow \\ \text{AR}(1) & \text{First} & & \text{MA}(1)\end{array}$$

Backshift notation for ARIMA

- ARIMA(p, d, q) model:

$$(1 - \phi_1 B - \dots - \phi_p B^p) \underset{\substack{\uparrow \\ \text{AR}(p)}}}{(1 - B)^d} y_t = c + (1 + \theta_1 B + \dots + \theta_q B^q) \underset{\substack{\uparrow \\ \text{MA}(q)}}{e_t}$$

d differences

Outline

1 Backshift notation reviewed

2 Seasonal ARIMA models

3 ARIMA vs ETS

4 Lab session 12

Seasonal ARIMA models

ARIMA	$\underbrace{(p, d, q)}$	$\underbrace{(P, D, Q)_m}$
	↑	↑
	Non-seasonal part of the model	Seasonal part of of the model

where m = number of observations per year.

Seasonal ARIMA models


E.g., $\text{ARIMA}(1, 1, 1)(1, 1, 1)_4$ model (without constant)

Seasonal ARIMA models

E.g., ARIMA(1, 1, 1)(1, 1, 1)₄ model (without constant)
$$(1 - \phi_1 B)(1 - \Phi_1 B^4)(1 - B)(1 - B^4)y_t = (1 + \theta_1 B)(1 + \Theta_1 B^4)e_t.$$

Seasonal ARIMA models

E.g., $\text{ARIMA}(1, 1, 1)(1, 1, 1)_4$ model (without constant)
 $(1 - \phi_1 B)(1 - \Phi_1 B^4)(1 - B)(1 - B^4)y_t = (1 + \theta_1 B)(1 + \Theta_1 B^4)e_t.$



Seasonal ARIMA models

E.g., ARIMA(1, 1, 1)(1, 1, 1)₄ model (without constant)
 $(1 - \phi_1 B)(1 - \Phi_1 B^4)(1 - B)(1 - B^4)y_t = (1 + \theta_1 B)(1 + \Theta_1 B^4)e_t.$

All the factors can be multiplied out and the general model written as follows:

$$\begin{aligned} y_t = & (1 + \phi_1)y_{t-1} - \phi_1 y_{t-2} + (1 + \Phi_1)y_{t-4} \\ & - (1 + \phi_1 + \Phi_1 + \phi_1 \Phi_1)y_{t-5} + (\phi_1 + \phi_1 \Phi_1)y_{t-6} \\ & - \Phi_1 y_{t-8} + (\Phi_1 + \phi_1 \Phi_1)y_{t-9} - \phi_1 \Phi_1 y_{t-10} \\ & + e_t + \theta_1 e_{t-1} + \Theta_1 e_{t-4} + \theta_1 \Theta_1 e_{t-5}. \end{aligned}$$

Common ARIMA models

In the US Census Bureau uses the following models most often:

ARIMA(0,1,1)(0,1,1)_m with log transformation

ARIMA(0,1,2)(0,1,1)_m with log transformation

ARIMA(2,1,0)(0,1,1)_m with log transformation

ARIMA(0,2,2)(0,1,1)_m with log transformation

ARIMA(2,1,2)(0,1,1)_m with no transformation

Seasonal ARIMA models

The seasonal part of an AR or MA model will be seen in the seasonal lags of the PACF and ACF.

$\text{ARIMA}(0,0,0)(0,0,1)_{12}$ will show:


- a spike at lag 12 in the ACF but no other significant spikes.
- The PACF will show exponential decay in the seasonal lags; that is, at lags 12, 24, 36,

$\text{ARIMA}(0,0,0)(1,0,0)_{12}$ will show:

- exponential decay in the seasonal lags of the ACF
- a single significant spike at lag 12 in the PACF.


European quarterly retail trade

```
autoplot(euretail) +  
  xlab("Year") + ylab("Retail index")
```




European quarterly retail trade

```
euretail %>% diff(lag=4) %>% ggtsdisplay()
```



European quarterly retail trade

```
euretail %>% diff(lag=4) %>% diff() %>%  
ggtsdisplay()
```




European quarterly retail trade

- $d = 1$ and $D = 1$ seems necessary.
- Significant spike at lag 1 in ACF suggests non-seasonal MA(1) component.
- Significant spike at lag 4 in ACF suggests seasonal MA(1) component.
- Initial candidate model: $\text{ARIMA}(0,1,1)(0,1,1)_4$.
- We could also have started with $\text{ARIMA}(1,1,0)(1,1,0)_4$.

European quarterly retail trade

```
fit <- Arima(euretail, order=c(0,1,1),  
            seasonal=c(0,1,1))  
checkresiduals(fit)
```

Residuals from ARIMA(0,1,1)(0,1,1)[4]



European quarterly retail trade

```
##
```

```
##  Ljung-Box test
```

```
##
```

```
## data:  Residuals from ARIMA(0,1,1)(0,1,1)[4
```

```
## Q* = 10.654, df = 6, p-value = 0.09968
```

```
##
```

```
## Model df: 2.    Total lags used: 8
```


European quarterly retail trade

- ACF and PACF of residuals show significant spikes at lag 2, and maybe lag 3.
- AICc of $\text{ARIMA}(0,1,2)(0,1,1)_4$ model is 74.36.
- AICc of $\text{ARIMA}(0,1,3)(0,1,1)_4$ model is 68.53.

European quarterly retail trade

- ACF and PACF of residuals show significant spikes at lag 2, and maybe lag 3.
- AICc of $\text{ARIMA}(0,1,2)(0,1,1)_4$ model is 74.36.
- AICc of $\text{ARIMA}(0,1,3)(0,1,1)_4$ model is 68.53.


```
fit <- Arima(euretail, order=c(0,1,3),  
            seasonal=c(0,1,1))  
checkresiduals(fit)
```

European quarterly retail trade

```
## Series: euretail
## ARIMA(0,1,3)(0,1,1)[4]
##
## Coefficients:
##          ma1      ma2      ma3      sma1
##      0.2630  0.3694  0.4200  -0.6636
## s.e.  0.1237  0.1255  0.1294   0.1545
##
## sigma^2 estimated as 0.156:  log likelihood=-
28.63
## AIC=67.26   AICc=68.39   BIC=77.65
```

European quarterly retail trade

`checkresiduals(fit)`




European quarterly retail trade

```
##  
##  Ljung-Box test  
##  
## data:  Residuals from ARIMA(0,1,3)(0,1,1)[4]  
## Q* = 0.51128, df = 4, p-value = 0.9724  
##  
## Model df: 4.    Total lags used: 8
```

European quarterly retail trade

```
autoplot(forecast(fit, h=12))
```



European quarterly retail trade

```
auto.arima(euretail)
```


```
## Series: euretail
## ARIMA(1,1,2)(0,1,1)[4]
##
## Coefficients:
##          ar1          ma1          ma2          sma1
##      0.7362   -0.4663   0.2163   -0.8433
## s.e.  0.2243    0.1990   0.2101    0.1876
##
## sigma^2 estimated as 0.1587:  log likelihood=-
29.62
## AIC=69.24    AICc=70.38    BIC=79.63
```

European quarterly retail trade


```
auto.arima(euretail, stepwise=FALSE, approximation=FALSE)
```

```
## Series: euretail
## ARIMA(0,1,3)(0,1,1)[4]
##
## Coefficients:
##           ma1      ma2      ma3      sma1
##      0.2630  0.3694  0.4200 -0.6636
## s.e.  0.1237  0.1255  0.1294  0.1545
##
## sigma^2 estimated as 0.156:  log likelihood=-
28.63
## AIC=67.26   AICc=68.39   BIC=77.65
```


Corticosteroid drug sales



Corticosteroid drug sales



Corticosteroid drug sales

- Choose $D = 1$ and $d = 0$.
- Spikes in PACF at lags 12 and 24 suggest seasonal AR(2) term.
- Spikes in PACF suggests possible non-seasonal AR(3) term.
- Initial candidate model: $\text{ARIMA}(3,0,0)(2,1,0)_{12}$.

Corticosteroid drug sales

Model	AICc
ARIMA(3,0,0)(2,1,0) ₁₂	-475.12
ARIMA(3,0,1)(2,1,0) ₁₂	-476.31
ARIMA(3,0,2)(2,1,0) ₁₂	-474.88
ARIMA(3,0,1)(1,1,0) ₁₂	-463.40
ARIMA(3,0,1)(0,1,1) ₁₂	-483.67
ARIMA(3,0,1)(0,1,2) ₁₂	-485.48
ARIMA(3,0,1)(1,1,1) ₁₂	-484.25

Corticosteroid drug sales

```
(fit <- Arima(h02, order=c(3,0,1), seasonal=c(0,1,2),  
  lambda=0))
```

```
## Series: h02
```

```
## ARIMA(3,0,1)(0,1,2)[12]
```

```
## Box Cox transformation: lambda= 0
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ar2      ar3      ma1      sma1      sma2
```

```
##      -0.1603  0.5481  0.5678  0.3827  -0.5222  -
```

```
0.1768
```

```
## s.e.    0.1636  0.0878  0.0942  0.1895  0.0861  0.0872
```


```
##
```

```
## sigma^2 estimated as 0.004278:  log likelihood=250.04 30
```

```
## AIC= 486.08  AICc= 485.48  BIC= 463.28
```

Corticosteroid drug sales

`checkresiduals(fit)`



Corticosteroid drug sales

##

Ljung-Box test

##

data: Residuals from ARIMA(3,0,1)(0,1,2)[1

Q* = 23.663, df = 18, p-value = 0.1664

##

Model df: 6. Total lags used: 24

Corticosteroid drug sales

```
(fit <- auto.arima(h02, lambda=0, d=0, D=1, max.order=9,  
  stepwise=FALSE, approximation=FALSE))
```

```
## Series: h02
```

```
## ARIMA(3,0,1)(0,1,2)[12] with drift
```

```
## Box Cox transformation: lambda= 0
```

```
##
```

```
## Coefficients:
```

```
##          ar1      ar2      ar3      ma1      sma1      sma2      drift
```

```
##          -0.2653  0.5011  0.5394  0.4572 -0.5031  -
```

```
0.2030  0.0038
```

```
## s.e.      0.1691  0.0813  0.0848  0.1904   0.0847   0.0871  0.0009
```

```
##
```

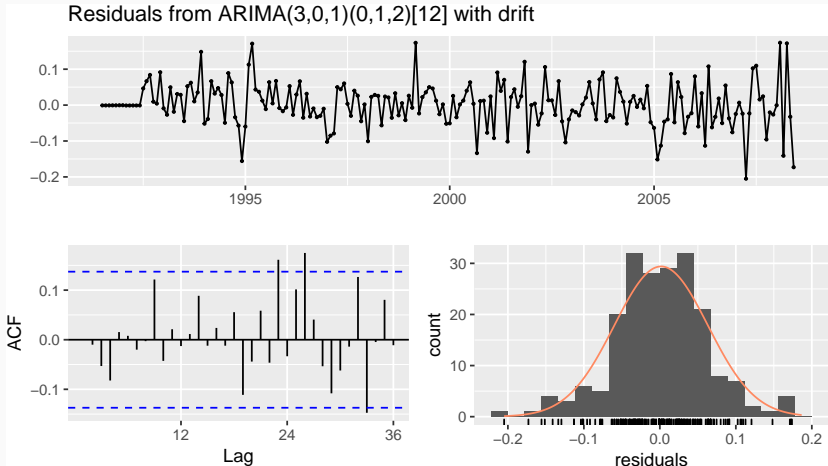
```
## sigma^2 estimated as 0.004176: log likelihood=252.99
```

33

```
## AIC= 480.00   AICc= 480.2   BIC= 462.02
```


Corticosteroid drug sales

`checkresiduals(fit)`



Corticosteroid drug sales

```
##
```

```
##  Ljung-Box test
```

```
##
```

```
## data:  Residuals from ARIMA(3,0,1)(0,1,2)[12] wi
```

```
## Q* = 19.369, df = 17, p-value = 0.3078
```

```
##
```

```
## Model df: 7.    Total lags used: 24
```

Corticosteroid drug sales

Training data: July 1991 to June 2006

Test data: July 2006–June 2008

```
getrmse <- function(x,h,...)
{
  train.end <- time(x)[length(x)-h]
  test.start <- time(x)[length(x)-h+1]
  train <- window(x,end=train.end)
  test <- window(x,start=test.start)
  fit <- Arima(train,...)
  fc <- forecast(fit,h=h)
  return(accuracy(fc,test)[2,"RMSE"])
}
```

```
getrmse(h02,h=24,order=c(3,0,0),seasonal=c(2,1,0),lambda=0)
```

```
getrmse(h02,h=24,order=c(3,0,1),seasonal=c(2,1,0),lambda=0)
```

```
getrmse(h02,h=24,order=c(3,0,2),seasonal=c(2,1,0),lambda=0)
```

```
getrmse(h02,h=24,order=c(2,0,1),seasonal=c(1,1,0),lambda=0)
```

Corticosteroid drug sales


Model	RMSE
ARIMA(3,0,0)(2,1,0)[12]	0.0661
ARIMA(3,0,1)(2,1,0)[12]	0.0646
ARIMA(3,0,2)(2,1,0)[12]	0.0645
ARIMA(3,0,1)(1,1,0)[12]	0.0679
ARIMA(3,0,1)(0,1,1)[12]	0.0644
ARIMA(3,0,1)(0,1,2)[12]	0.0622
ARIMA(3,0,1)(1,1,1)[12]	0.0630
ARIMA(4,0,3)(0,1,1)[12]	0.0648
ARIMA(3,0,3)(0,1,1)[12]	0.0639
ARIMA(4,0,2)(0,1,1)[12]	0.0648
ARIMA(3,0,2)(0,1,1)[12]	0.0644
ARIMA(2,1,3)(0,1,1)[12]	0.0634
ARIMA(2,1,4)(0,1,1)[12]	0.0632

Corticosteroid drug sales

- Models with lowest AICc values tend to give slightly better results than the other models.
- AICc comparisons must have the same orders of differencing. But RMSE test set comparisons can involve any models.
- No model passes all the residual tests.
- Use the best model available, even if it does not pass all tests.
- In this case, the $\text{ARIMA}(3,0,1)(0,1,2)_{12}$ has the lowest RMSE value and the best AICc value for models with fewer than 6 parameters.

Corticosteroid drug sales

```
fit <- Arima(h02, order=c(3,0,1), seasonal=c(0,1,2),  
  lambda=0)  
autoplot(forecast(fit)) +  
  ylab("H02 sales (million scripts)") + xlab("Year")
```



Outline

1 Backshift notation reviewed

2 Seasonal ARIMA models

3 ARIMA vs ETS

4 Lab session 12

ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.
- Many ARIMA models have no exponential smoothing counterparts.
- ETS models all non-stationary. Models with seasonality or non-damped trend (or both) have

Equivalences

ETS model	ARIMA model	Parameters
ETS(A,N,N)	ARIMA(0,1,1)	$\theta_1 = \alpha - 1$
ETS(A,A,N)	ARIMA(0,2,2)	$\theta_1 = \alpha + \beta - 2$ $\theta_2 = 1 - \alpha$
ETS(A,A,N)	ARIMA(1,1,2)	$\phi_1 = \phi$ $\theta_1 = \alpha + \phi\beta - 1 - \phi$ $\theta_2 = (1 - \alpha)\phi$
ETS(A,N,A)	ARIMA(0,0,m)(0,1,0) _m	
ETS(A,A,A)	ARIMA(0,1,m + 1)(0,1,0) _m	
ETS(A,A,A)	ARIMA(1,0,m + 1)(0,1,0) _m	

Outline

1 Backshift notation reviewed

2 Seasonal ARIMA models

3 ARIMA vs ETS

4 Lab session 12

Lab Session 12